

CO527: Advanced Database Systems

The Relational Algebra

Sampath Deegalla
Senior Lecturer, Department of Computer Engineering
University of Peradeniya
dsdeegalla@pdn.ac.lk

11th February 2020

- Relational algebra is a query language that allows us to retrieve data from DBs.
- It consists of a set of operations that take one or two relations as input and produce a new relation as output.
- The result of a retrieval is a new relation, which may have been formed from one or more relations. The **algebra operations** thus produce new relations, which can be further manipulated using operations of the same algebra.
- A sequence of relational algebra operations forms a **relational algebra expression**, whose result will also be a relation that represents the result of a database query (or retrieval request).

- Relational algebra is a query language that allows us to retrieve data from DBs.
- It consists of a set of operations that take one or two relations as input and produce a new relation as output.
- The result of a retrieval is a new relation, which may have been formed from one or more relations. The **algebra operations** thus produce new relations, which can be further manipulated using operations of the same algebra.
- A sequence of relational algebra operations forms a **relational algebra expression**, whose result will also be a relation that represents the result of a database query (or retrieval request).

- Relational algebra is a query language that allows us to retrieve data from DBs.
- It consists of a set of operations that take one or two relations as input and produce a new relation as output.
- The result of a retrieval is a new relation, which may have been formed from one or more relations. The **algebra operations** thus produce new relations, which can be further manipulated using operations of the same algebra.
- A sequence of relational algebra operations forms a **relational algebra expression**, whose result will also be a relation that represents the result of a database query (or retrieval request).

- Relational algebra is a query language that allows us to retrieve data from DBs.
- It consists of a set of operations that take one or two relations as input and produce a new relation as output.
- The result of a retrieval is a new relation, which may have been formed from one or more relations. The **algebra operations** thus produce new relations, which can be further manipulated using operations of the same algebra.
- A sequence of relational algebra operations forms a **relational algebra expression**, whose result will also be a relation that represents the result of a database query (or retrieval request).

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

- The SELECT operation is used to select a subset of the tuples from a relation that satisfy a **selection condition**.
- It is a filter that keeps only those tuples that satisfy a qualifying condition.
- Those satisfying the condition are selected while others are not included in the result.
- In general, the select operation is denoted by

- The SELECT operation is used to select a subset of the tuples from a relation that satisfy a **selection condition**.
- It is a filter that keeps only those tuples that satisfy a qualifying condition.
- Those satisfying the condition are selected while others are not included in the result.
- In general, the select operation is denoted by

- The SELECT operation is used to select a subset of the tuples from a relation that satisfy a **selection condition**.
- It is a filter that keeps only those tuples that satisfy a qualifying condition.
- Those satisfying the condition are selected while others are not included in the result.
- In general, the select operation is denoted by

- $$\sigma_{\langle \text{selection condition} \rangle}(R)$$

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ 🔍 ↺

- $$\sigma_{\text{DNO}=4}(\text{EMPLOYEE})$$

- $$\sigma_{\text{DNO}=4}(\textit{EMPLOYEE})$$

- $$\sigma_{\text{DNO}=4}(\text{EMPLOYEE})$$

- $$\sigma_{\text{SALARY} > 30,000}(\text{EMPLOYEE})$$

- A selection condition is a set of clauses connected by the Boolean operators AND, OR, and NOT
- Each clause has the format
 <attribute name> <comparison op> < constant value>
or
 <attribute name> <comparison op> <attribute name>
- The comparison operators are =, <, ≤, >, ≥ and ≠

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is **commutative**; i.e.,

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is commutative; i.e.,

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is **commutative**; i.e.,

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is **commutative**; i.e.,

$$\sigma_{\langle \text{cond } 1 \rangle}(\sigma_{\langle \text{cond } 2 \rangle}(R)) = \sigma_{\langle \text{cond } 2 \rangle}(\sigma_{\langle \text{cond } 1 \rangle}(R))$$

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is **commutative**; i.e.,

$$\sigma_{\langle \text{cond } 1 \rangle}(\sigma_{\langle \text{cond } 2 \rangle}(R)) = \sigma_{\langle \text{cond } 2 \rangle}(\sigma_{\langle \text{cond } 1 \rangle}(R))$$

- A cascaded SELECT operation may be replaced by a single selection with a conjunction of all the conditions; i.e.,

- SELECT is a unary operator that takes one relation as input and produces another relation as output
- The SELECT operation $\sigma_{\langle \text{selection condition} \rangle}(R)$ produces a relation S that has the same schema as R
- The SELECT operation σ is **commutative**; i.e.,

$$\sigma_{\langle \text{cond } 1 \rangle}(\sigma_{\langle \text{cond } 2 \rangle}(R)) = \sigma_{\langle \text{cond } 2 \rangle}(\sigma_{\langle \text{cond } 1 \rangle}(R))$$

- A cascaded SELECT operation may be replaced by a single selection with a conjunction of all the conditions; i.e.,

$$\sigma_{\langle \text{cond } 1 \rangle}(\sigma_{\langle \text{cond } 2 \rangle}(\sigma_{\langle \text{cond } 3 \rangle}(R))) = \sigma_{\langle \text{cond } 1 \rangle \text{ AND } \langle \text{cond } 2 \rangle \text{ AND } \langle \text{cond } 3 \rangle}(R)$$

Select Example

To select the tuples for all employees who either work in department 4 and make \$25,000 per year, or work in department 5 and make over \$30,000:

Select Example

To select the tuples for all employees who either work in department 4 and make \$25,000 per year, or work in department 5 and make over \$30,000:

$$(a) \sigma_{(DNO=4 \text{ AND } SALARY > 25000) \text{ OR } (DNO=5 \text{ AND } SALARY > 30000)}(EMPLOYEE)$$

Select Example

To select the tuples for all employees who either work in department 4 and make \$25,000 per year, or work in department 5 and make over \$30,000:

(a) $\sigma_{(DNO=4 \text{ AND } SALARY > 25000) \text{ OR } (DNO=5 \text{ AND } SALARY > 30000)}(EMPLOYEE)$

(a)

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
Franklin	T	Wong	333445555	1955-12-08	638 Voss,Houston,TX	M	40000	888665555	5
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry,Bellaire,TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 FireOak,Humble,TX	M	38000	333445555	5

- $$\pi_{\langle \text{attribute list} \rangle}(R)$$

- The project operation removes any duplicate tuples, so the result of the project operation is a set of tuples and hence a valid relation.

- $$\pi_{\text{LNAME, FNAME, SALARY}}(\text{EMPLOYEE})$$

- The number of tuples in the result of projection $\pi_{\langle \text{list} \rangle}(R)$ is always less or equal to the number of tuples in R .
- If the list of attributes includes a key of R , then the number of tuples is equal to the number of tuples in R .
- $\pi_{\langle \text{list1} \rangle}(\pi_{\langle \text{list2} \rangle}(R)) = \pi_{\langle \text{list1} \rangle}(R)$ as long as $\langle \text{list2} \rangle$ contains the attributes in $\langle \text{list1} \rangle$

(c) $\pi_{\text{SEX, SALARY}}(\text{EMPLOYEE})$

Project Example

(b) $\pi_{\text{LNAME, FNAME, SALARY}}(\text{EMPLOYEE})$

(c) $\pi_{\text{SEX, SALARY}}(\text{EMPLOYEE})$

(b)

LNAME	FNAME	SALARY
Smith	John	30000
Wong	Franklin	40000
Zelaya	Alicia	25000
Wallace	Jennifer	43000
Narayan	Ramesh	38000
English	Joyce	25000
Jabbar	Ahmad	25000
Borg	James	55000

(c)

SEX	SALARY
M	30000
M	40000
F	25000
F	43000
M	38000
M	25000
M	55000

Renaming of Relational Operations

- We may want to apply several relational algebra operations one after the other.
- Either we can write the operations as a single relational algebra expression by nesting the operations, or we can apply one operation at a time and create intermediate result relations. In the latter case, we must give names to the relations that hold the intermediate results.

Renaming of Relational Operations

- Example: To retrieve the first name, last name, and salary of all employees who work in department number 5, we must apply a select and a project operation. We can write a single relational algebra expression as follows:

$$\pi_{\text{FNAME, LNAME, SALARY}}(\sigma_{\text{DNO} = 5}(\text{EMPLOYEE}))$$

- OR We can explicitly show the sequence of operations, giving a name to each intermediate relation:

$$\text{DEP5_EMPS} \leftarrow \sigma_{\text{DNO} = 5}(\text{EMPLOYEE})$$

$$\text{RESULT} \leftarrow \pi_{\text{FNAME, LNAME, SALARY}}(\text{DEP5_EMPS})$$

- We can also rename the attributes, if desired, by specifying new attribute names when we name a partial result, e.g., RESULT (F, L, S)

(b) $TEMP \leftarrow \sigma_{DNO=5}(EMPLOYEE)$

$$R(\text{FIRSTNAME}, \text{LASTNAME}, \text{SALARY}) \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SALARY}}(TEMP)$$

(b)	TEMP	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
		John	B	Smith	123456789	1965-01-09	731 Fondren,Houston,TX	M	30000	333445555	5
		Franklin	T	Wong	333445555	1955-12-08	638 Voss,Houston,TX	M	40000	888665555	5
		Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak,Humble,TX	M	38000	333445555	5
		Joyce	A	English	453453453	1972-07-31	5631 Rice,Houston,TX	F	25000	333445555	5

R	FIRSTNAME	LASTNAME	SALARY
	John	Smith	30000
	Franklin	Wong	40000
	Ramesh	Narayan	38000
	Joyce	English	25000

Exercise

- 1 Which projects are located in Houston?
- 2 What are the names of the departments?
- 3 Find out everything about all of the employees who were born before 1950-01-01.
- 4 What are the names of the employees who were born before 1950-01-01?
- 5 When did the manager of the Research department begin managing that department?

○○○○○○○○○○○○○●○○○○○○○○○○○○○○○○○○○○○○○○○○○○○○

[illegible]

1999-01-01	0	1
------------	---	---

[illegible]

Year	Number of cases	Rate per 100,000
1990	1,000	1.0
1991	1,100	1.1
1992	1,200	1.2
1993	1,300	1.3
1994	1,400	1.4
1995	1,500	1.5
1996	1,600	1.6
1997	1,700	1.7
1998	1,800	1.8
1999	1,900	1.9
2000	2,000	2.0
2001	2,100	2.1
2002	2,200	2.2
2003	2,300	2.3
2004	2,400	2.4
2005	2,500	2.5
2006	2,600	2.6
2007	2,700	2.7
2008	2,800	2.8
2009	2,900	2.9
2010	3,000	3.0
2011	3,100	3.1
2012	3,200	3.2
2013	3,300	3.3
2014	3,400	3.4
2015	3,500	3.5
2016	3,600	3.6
2017	3,700	3.7
2018	3,800	3.8
2019	3,900	3.9
2020	4,000	4.0

F. J. Beckwith, Editor	99	Oxford	4
------------------------	----	--------	---

1990-1991	1991-1992	1992-1993	1993-1994	1994-1995	1995-1996	1996-1997	1997-1998	1998-1999	1999-2000	2000-2001	2001-2002	2002-2003	2003-2004	2004-2005	2005-2006	2006-2007	2007-2008	2008-2009	2009-2010	2010-2011	2011-2012	2012-2013	2013-2014	2014-2015	2015-2016	2016-2017	2017-2018	2018-2019	2019-2020	2020-2021	2021-2022	2022-2023	2023-2024	2024-2025	2025-2026	2026-2027	2027-2028	2028-2029	2029-2030	2030-2031	2031-2032	2032-2033	2033-2034	2034-2035	2035-2036	2036-2037	2037-2038	2038-2039	2039-2040	2040-2041	2041-2042	2042-2043	2043-2044	2044-2045	2045-2046	2046-2047	2047-2048	2048-2049	2049-2050	2050-2051	2051-2052	2052-2053	2053-2054	2054-2055	2055-2056	2056-2057	2057-2058	2058-2059	2059-2060	2060-2061	2061-2062	2062-2063	2063-2064	2064-2065	2065-2066	2066-2067	2067-2068	2068-2069	2069-2070	2070-2071	2071-2072	2072-2073	2073-2074	2074-2075	2075-2076	2076-2077	2077-2078	2078-2079	2079-2080	2080-2081	2081-2082	2082-2083	2083-2084	2084-2085	2085-2086	2086-2087	2087-2088	2088-2089	2089-2090	2090-2091	2091-2092	2092-2093	2093-2094	2094-2095	2095-2096	2096-2097	2097-2098	2098-2099	2099-2100	2100-2101	2101-2102	2102-2103	2103-2104	2104-2105	2105-2106	2106-2107	2107-2108	2108-2109	2109-2110	2110-2111	2111-2112	2112-2113	2113-2114	2114-2115	2115-2116	2116-2117	2117-2118	2118-2119	2119-2120	2120-2121	2121-2122	2122-2123	2123-2124	2124-2125	2125-2126	2126-2127	2127-2128	2128-2129	2129-2130	2130-2131	2131-2132	2132-2133	2133-2134	2134-2135	2135-2136	2136-2137	2137-2138	2138-2139	2139-2140	2140-2141	2141-2142	2142-2143	2143-2144	2144-2145	2145-2146	2146-2147	2147-2148	2148-2149	2149-2150	2150-2151	2151-2152	2152-2153	2153-2154	2154-2155	2155-2156	2156-2157	2157-2158	2158-2159	2159-2160	2160-2161	2161-2162	2162-2163	2163-2164	2164-2165	2165-2166	2166-2167	2167-2168	2168-2169	2169-2170	2170-2171	2171-2172	2172-2173	2173-2174	2174-2175	2175-2176	2176-2177	2177-2178	2178-2179	2179-2180	2180-2181	2181-2182	2182-2183	2183-2184	2184-2185	2185-2186	2186-2187	2187-2188	2188-2189	2189-2190	2190-2191	2191-2192	2192-2193	2193-2194	2194-2195	2195-2196	2196-2197	2197-2198	2198-2199	2199-2200	2200-2201	2201-2202	2202-2203	2203-2204	2204-2205	2205-2206	2206-2207	2207-2208	2208-2209	2209-2210	2210-2211	2211-2212	2212-2213	2213-2214	2214-2215	2215-2216	2216-2217	2217-2218	2218-2219	2219-2220	2220-2221	2221-2222	2222-2223	2223-2224	2224-2225	2225-2226	2226-2227	2227-2228	2228-2229	2229-2230	2230-2231	2231-2232	2232-2233	2233-2234	2234-2235	2235-2236	2236-2237	2237-2238	2238-2239	2239-2240	2240-2241	2241-2242	2242-2243	2243-2244	2244-2245	2245-2246	2246-2247	2247-2248	2248-2249	2249-2250	2250-2251	2251-2252	2252-2253	2253-2254	2254-2255	2255-2256	2256-2257	2257-2258	2258-2259	2259-2260	2260-2261	2261-2262	2262-
-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-------

Set Theoretic Relational Operators

- The set theoretic operators are **union** ($R \cup S$), **intersection** ($R \cap S$) and **difference** ($R - S$).
- Since relations are sets of tuples, we can borrow established operators that work on sets.
- These are all **binary operators**. They each take two relations as operands and produce one relation as their result.
- They all require that their input relations are **union compatible**.

Union Compatibility

- For two operand relations, $R(A_1, A_2, \dots, A_n)$ and $S(B_1, B_2, \dots, B_n)$ to be union compatible,
 - they must have the **same number of attributes**, and
 - the **domains of their corresponding attributes must be the same**; that is, $dom(A_i) = dom(B_i)$ for $i = 1, 2, \dots, n$.
- The relation that results from a set theoretic operation will also be union compatible with the two input relations.
- The names of the corresponding attributes do not have to be the same.

The Union Operation

- The result of the UNION operation, denoted by $R \cup S$, is a relation that includes **all tuples that are either in R or in S or in both R and S**. Duplicate tuples are eliminated.
 - Example: To retrieve the social security numbers of all employees who either work in department 5 or directly supervise an employee who works in department 5, we can use the union operation as follows:

$$DEP5_EMPS \leftarrow \sigma_{DNO=5}(EMPLOYEE)$$

$$RESULT1 \leftarrow \pi_{SSN}(DEP5_EMPS)$$

$$RESULT2(SSN) \leftarrow \pi_{SUPERSSN}(DEP5_EMPS)$$

$$RESULT \leftarrow RESULT1 \cup RESULT2$$

- The union operation produces the tuples that are in either RESULT1 or RESULT2 or both.

The Union Operation

- The result of the UNION operation, denoted by $R \cup S$, is a relation that includes **all tuples that are either in R or in S or in both R and S**. Duplicate tuples are eliminated.
 - Example: To retrieve the social security numbers of all employees who either work in department 5 or directly supervise an employee who works in department 5, we can use the union operation as follows:

$$DEP5_EMPS \leftarrow \sigma_{DNO = 5}(EMPLOYEE)$$

$$RESULT1 \leftarrow \pi_{SSN}(DEP5_EMPS)$$

$$RESULT2(SSN) \leftarrow \pi_{SUPERSSN}(DEP5_EMPS)$$

$$RESULT \leftarrow RESULT1 \cup RESULT2$$

- The union operation produces the tuples that are in either RESULT1 or RESULT2 or both.

The Union Operation

- The result of the UNION operation, denoted by $R \cup S$, is a relation that includes **all tuples that are either in R or in S or in both R and S**. Duplicate tuples are eliminated.
 - Example: To retrieve the social security numbers of all employees who either work in department 5 or directly supervise an employee who works in department 5, we can use the union operation as follows:

$$DEP5_EMPS \leftarrow \sigma_{DNO=5}(EMPLOYEE)$$

$$RESULT1 \leftarrow \pi_{SSN}(DEP5_EMPS)$$

$$RESULT2(SSN) \leftarrow \pi_{SUPERSSN}(DEP5_EMPS)$$

$$RESULT \leftarrow RESULT1 \cup RESULT2$$

- The union operation produces the tuples that are in either RESULT1 or RESULT2 or both.

The Union Operation

- The result of the UNION operation, denoted by $R \cup S$, is a relation that includes **all tuples that are either in R or in S or in both R and S**. Duplicate tuples are eliminated.
 - Example: To retrieve the social security numbers of all employees who either work in department 5 or directly supervise an employee who works in department 5, we can use the union operation as follows:

$$DEP5_EMPS \leftarrow \sigma_{DNO=5}(EMPLOYEE)$$

$$RESULT1 \leftarrow \pi_{SSN}(DEP5_EMPS)$$

$$RESULT2(SSN) \leftarrow \pi_{SUPERSSN}(DEP5_EMPS)$$

$$RESULT \leftarrow RESULT1 \cup RESULT2$$

- The union operation produces the tuples that are in either RESULT1 or RESULT2 or both.

The Intersection Operation

- The result of the INTERSECTION operation, denoted by $R \cap S$, is a relation that includes **all tuples that are in both R and S**. The two operands must be union compatible.
 - Example: The result of the intersection operation includes only those who are both students and instructors.
 - $STUDENT \cap INSTRUCTOR$

(a)

STUDENT	FN	LN
	Susan	Yao
	Ramesh	Shah
	Johnny	Kohler
	Barbara	Jones
	Amy	Ford
	Jimmy	Wang
	Ernest	Gilbert

INSTRUCTOR	FNAME	LNAME
	John	Smith
	Ricardo	Browne
	Susan	Yao
	Francis	Johnson
	Ramesh	Shah

(c)

FN	LN
Susan	Yao
Ramesh	Shah

The Set Difference Operation

- The result of the SET DIFFERENCE operation, denoted by $R - S$, is a relation that includes **all tuples that are in R but not in S**. The two operands must be union compatible.
 - Example: The figure shows the names of students who are not instructors, and the names of instructors who are not students.
 - (d) STUDENT - INSTRUCTOR (e) INSTRUCTOR - STUDENT

(a)

STUDENT	FN	LN
	Susan	Yao
	Ramesh	Shah
	Johnny	Kohler
	Barbara	Jones
	Amy	Ford
	Jimmy	Wang
	Ernest	Gilbert

INSTRUCTOR	FNAME	LNAME
	John	Smith
	Ricardo	Browne
	Susan	Yao
	Francis	Johnson
	Ramesh	Shah

(d)

FN	LN
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

(e)

FNAME	LNAME
John	Smith
Ricardo	Browne
Francis	Johnson

Examples on UNION, INTERSECTION and MINUS

(b) STUDENT \cup INSTRUCTOR. (c) STUDENT \cap INSTRUCTOR. (d) STUDENT – INSTRUCTOR. (e) INSTRUCTOR – STUDENT

(a)

STUDENT	FN	LN
	Susan	Yao
	Ramesh	Shah
	Johnny	Kohler
	Barbara	Jones
	Amy	Ford
	Jimmy	Wang
	Ernest	Gilbert

INSTRUCTOR	FNAME	LNAME
	John	Smith
	Ricardo	Browne
	Susan	Yao
	Francis	Johnson
	Ramesh	Shah

(b)

FN	LN
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert
John	Smith
Ricardo	Browne
Francis	Johnson

(c)

FN	LN
Susan	Yao
Ramesh	Shah

(d)

FN	LN
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

(e)

FNAME	LNAME
John	Smith
Ricardo	Browne
Francis	Johnson

Set Theoretic Relational Operators

- Notice that both union and intersection are **commutative** operations; that is

$$R \cup S = S \cup R, \text{ and } R \cap S = S \cap R$$

- Both union and intersection can be treated as n-ary operations applicable to any number of relations as both are **associative** operations; that is

$$R \cup (S \cup T) = (R \cup S) \cup T, \text{ and}$$

$$(R \cap S) \cap T = R \cap (S \cap T)$$

- The minus operation is **not commutative**; that is, in general

$$R - S \neq S - R$$

Exercise

- 1 Find the social security numbers of the managers of departments who work on project 30.
- 2 Find the social security numbers of all employees who have dependents or who work on project 2.
- 3 Find the social security numbers of all employees who do not have any dependents.

The Cartesian Product (Cross Product)

- The CARTESION PRODUCT operation (also known as Cross Product or Cross Join) is used to combine tuples from two relations in a combinatorial fashion. In general, the result of

$$R(A_1, A_2, \dots, A_n) \times S(B_1, B_2, \dots, B_m)$$

is a relation Q with degree $n + m$ attributes

$$Q(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_m)$$

in that order.

- The resulting relation Q has one tuple for each combination of tuples – one from R and one from S .
- Hence, if R has n_R tuples (denoted as $|R| = n_R$), and S has n_S tuples, then $|R \times S|$ will have $n_R * n_S$ tuples.
- The two operands do NOT have to be union compatible

Cartesian Product Example

Example: To retrieve a list of names of each female employee's dependents

$$FEMALE_EMPS \leftarrow \sigma_{SEX = 'F'}(EMPLOYEE)$$
$$EMP_NAMES \leftarrow \pi_{FNAME, LNAME, SSN}(FEMALE_EMPS)$$
$$EMP_DEPENDENTS \leftarrow EMP_NAMES \times DEPENDENT$$

Cartesian Product Example

Example: To retrieve a list of names of each female employee's dependents

$$FEMALE_EMPS \leftarrow \sigma_{SEX = 'F'}(EMPLOYEE)$$

$$EMP_NAMES \leftarrow \pi_{FNAME, LNAME, SSN}(FEMALE_EMPS)$$

$$EMP_DEPENDENTS \leftarrow EMP_NAMES \times DEPENDENT$$

Cartesian Product Example

Example: To retrieve a list of names of each female employee's dependents

$$\text{FEMALE_EMPS} \leftarrow \sigma_{\text{SEX} = 'F'}(\text{EMPLOYEE})$$

$$\text{EMP_NAMES} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SSN}}(\text{FEMALE_EMPS})$$

$$\text{EMP_DEPENDENTS} \leftarrow \text{EMP_NAMES} \times \text{DEPENDENT}$$

Cartesian Product Example

Example: To retrieve a list of names of each female employee's dependents

$$FEMALE_EMPS \leftarrow \sigma_{SEX = 'F'}(EMPLOYEE)$$

$$EMPNAMES \leftarrow \pi_{FNAME, LNAME, SSN}(FEMALE_EMPS)$$

$$EMP_DEPENDENTS \leftarrow EMPNAMES \times DEPENDENT$$

Example

FEMALE_EMPS	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5

EMPNAME	FNAME	LNAME	SSN
	Alicia	Zelaya	999887777
	Jennifer	Wallace	987654321
	Joyce	English	453453453

Result

EMP_DEPENDENTS	FNAME	LNAME	SSN	ESSN	DEPENDENT_NAME	SEX	BDATE	...
	Alicia	Zelaya	999887777	333445555	Alice	F	1986-04-05	...
	Alicia	Zelaya	999887777	333445555	Theodore	M	1983-10-25	...
	Alicia	Zelaya	999887777	333445555	Joy	F	1958-05-03	...
	Alicia	Zelaya	999887777	987654321	Abner	M	1942-02-28	...
	Alicia	Zelaya	999887777	123456789	Michael	M	1988-01-04	...
	Alicia	Zelaya	999887777	123456789	Alice	F	1988-12-30	...
	Alicia	Zelaya	999887777	123456789	Elizabeth	F	1967-05-05	...
	Jennifer	Wallace	987654321	333445555	Alice	F	1986-04-05	...
	Jennifer	Wallace	987654321	333445555	Theodore	M	1983-10-25	...
	Jennifer	Wallace	987654321	333445555	Joy	F	1958-05-03	...
	Jennifer	Wallace	987654321	987654321	Abner	M	1942-02-28	...
	Jennifer	Wallace	987654321	123456789	Michael	M	1988-01-04	...
	Jennifer	Wallace	987654321	123456789	Alice	F	1988-12-30	...
	Jennifer	Wallace	987654321	123456789	Elizabeth	F	1967-05-05	...
	Joyce	English	453453453	333445555	Alice	F	1986-04-05	...
	Joyce	English	453453453	333445555	Theodore	M	1983-10-25	...
	Joyce	English	453453453	333445555	Joy	F	1958-05-03	...
	Joyce	English	453453453	987654321	Abner	M	1942-02-28	...
	Joyce	English	453453453	123456789	Michael	M	1988-01-04	...
	Joyce	English	453453453	123456789	Alice	F	1988-12-30	...
	Joyce	English	453453453	123456789	Elizabeth	F	1967-05-05	...

Result

EMP_DEPENDENTS	FNAME	LNAME	SSN	ESSN	DEPENDENT_NAME	SEX	BDATE	• • •
	Alicia	Zelaya	999887777	333445555	Alice	F	1986-04-05	• • •
	Alicia	Zelaya	999887777	333445555	Theodore	M	1983-10-25	• • •
	Alicia	Zelaya	999887777	333445555	Joy	F	1958-05-03	• • •
	Alicia	Zelaya	999887777	987654321	Abner	M	1942-02-28	• • •
	Alicia	Zelaya	999887777	123456789	Michael	M	1988-01-04	• • •
	Alicia	Zelaya	999887777	123456789	Alice	F	1988-12-30	• • •
	Alicia	Zelaya	999887777	123456789	Elizabeth	F	1967-05-05	• • •
	Jennifer	Wallace	987654321	333445555	Alice	F	1986-04-05	• • •
	Jennifer	Wallace	987654321	333445555	Theodore	M	1983-10-25	• • •
	Jennifer	Wallace	987654321	333445555	Joy	F	1958-05-03	• • •
	Jennifer	Wallace	987654321	987654321	Abner	M	1942-02-28	• • •
	Jennifer	Wallace	987654321	123456789	Michael	M	1988-01-04	• • •
	Jennifer	Wallace	987654321	123456789	Alice	F	1988-12-30	• • •
	Jennifer	Wallace	987654321	123456789	Elizabeth	F	1967-05-05	• • •
	Joyce	English	453453453	333445555	Alice	F	1986-04-05	• • •
	Joyce	English	453453453	333445555	Theodore	M	1983-10-25	• • •
	Joyce	English	453453453	333445555	Joy	F	1958-05-03	• • •
	Joyce	English	453453453	987654321	Abner	M	1942-02-28	• • •
	Joyce	English	453453453	123456789	Michael	M	1988-01-04	• • •
	Joyce	English	453453453	123456789	Alice	F	1988-12-30	• • •
	Joyce	English	453453453	123456789	Elizabeth	F	1967-05-05	• • •

Result

$$\text{ACTUAL_DEPENDENTS} \leftarrow \sigma_{\text{SSN}=\text{ESSN}}(\text{EMP_DEPENDENTS})$$

$$\text{RESULT} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{DEPENDENT_NAME}}(\text{ACTUAL_DEPENDENTS})$$

ACTUAL_DEPENDENTS	FNAME	LNAME	SSN	ESSN	DEPENDENT_NAME	SEX	BDATE	...
	Jennifer	Wallace	987654321	987654321	Abner	M	1942-02-28	...

RESULT	FNAME	LNAME	DEPENDENT_NAME
	Jennifer	Wallace	Abner

Join Operations

- We don't usually want a Cartesian product per se. We usually want some subset of it.
- So, we typically need to use the select operation to get to the part of the Cartesian product we want.
- We use JOIN operations for this. There are three types of JOIN operations, the theta join, the equijoin, and the natural join.
- The general form of a join operation on two relations $R(A_1, A_2, \dots, A_n)$ and $S(B_1, B_2, \dots, B_m)$ is:

$$R \bowtie_{\langle \text{join condition} \rangle} S$$

The Theta Join

- The theta join produces all combinations of tuples from two relations, R and S, that satisfy a join condition
- A join condition is similar to a select condition, except that you can not use the Boolean OR and NOT operators. All clauses are ANDed together
- If you need a more general condition, you can use select with a Cartesian product

$$R \bowtie_{\langle \text{condition} \rangle} S = \sigma_{\langle \text{condition} \rangle} (R \times S)$$

Join Example

- Suppose that we want to retrieve the name of the manager of each department. To get the manager's name, we need to combine each DEPARTMENT tuple with the EMPLOYEE tuple whose SSN value matches the MGRSSN value in the department tuple. We do this by using the join \bowtie operation.

$DEPT_MGR \leftarrow DEPARTMENT \bowtie_{MGRSSN = SSN} EMPLOYEE$

$RESULT \leftarrow \pi_{DNAME, LNAME, FNAME}(DEPT_MGR)$

Join Example

- Suppose that we want to retrieve the name of the manager of each department. To get the manager's name, we need to combine each DEPARTMENT tuple with the EMPLOYEE tuple whose SSN value matches the MGRSSN value in the department tuple. We do this by using the join \bowtie operation.

$$DEPT_MGR \leftarrow DEPARTMENT \bowtie_{MGRSSN = SSN} EMPLOYEE$$

$$RESULT \leftarrow \pi_{DNAME, LNAME, FNAME}(DEPT_MGR)$$

Equi-joins

- The most common use of join involves join conditions with equality comparisons only. Such a join, where the only comparison operator used is $=$, is called an EQUIJOIN. In the result of an EQUIJOIN we always have one or more pairs of attributes (whose names need not be identical) that have identical values in every tuple.

Natural Joins

- Because one of each pair of attributes with identical values is superfluous, a new operation called natural join - denoted by $*$ - was created to get rid of the second (unnecessary) attribute in an EQUIJOIN condition.
- The standard definition of natural join requires that the two join attributes, or each pair of corresponding join attributes, have the same name in both relations. If this is not the case, a renaming operation is applied first.

Natural Join Example

- To apply a natural join on the DNUMBER attributes of DEPARTMENT and DEPT_LOCATIONS, we write:

$$DEPT_LOCS \leftarrow DEPARTMENT * DEPT_LOCATIONS$$

The Outer Join Operation

- In NATURAL JOIN tuples without a matching (or related) tuple are eliminated from the join result. Tuples with null in the join attributes are also eliminated. Sometimes, as a practical matter, we want to keep this information.
- A set of operations, called outer joins, can be used when we want to keep all the tuples in R, or all those in S, or all those in both relations in the result of the join, regardless of whether or not they have matching tuples in the other relation.

The Outer Join Operation

- The left outer join operation keeps every tuple in the first or left relation R in $R \sqcup \bowtie S$; if no matching tuple is found in S , then the attributes of S in the join result are filled or “padded” with null values.
- A similar operation, right outer join, keeps every tuple in the second or right relation S in the result of $R \bowtie \sqcup S$.
- A third operation, full outer join, denoted by $\sqcup \bowtie \sqcup$ keeps all tuples in both the left and the right relations when no matching tuples are found, padding them with null values as needed.

Outer Join Example

- To find the names of all employees, along with the names of the departments they manage, we could use:

$$TEMP \leftarrow EMPLOYEE \bowtie_{SSN = MGRSSN} DEPARTMENT$$

$$RESULT \leftarrow \pi_{FNAME, MINIT, LNAME, DNAME}(TEMP)$$

Complete Set of Relational Operations

- The set of operations including **select** σ , **project** π , **union** \cup , **set difference** $-$, and **Cartesian product** \times is called a **complete set** because any other relational algebra expression can be expressed by a combination of these five operations.
- For example:

$$R \cap S = (R \cup S) - ((R - S) \cup (S - R))$$

$$R \bowtie_{\langle \text{join condition} \rangle} S = \sigma_{\langle \text{join condition} \rangle} (R \times S)$$

Additional Relational Operations

- Even though we could, in principle, get by with only five relational algebra operations, in practice, we use more.
- Sometimes, these are convenient shortcuts, like the join operation.
- Sometimes, we want to extend the relational algebra to enable us to make some types of queries that were not originally supported.

The Division Operation

- DIVISION is another shortcut operation that could be expressed using only π , \times , and $-$
- It is applied to two relations, $R(Z) \div S(X)$, where the attributes X are a subset of the attributes Z . Let $Y = Z - X$ (and hence $Z = X \cup Y$); that is, let Y be the set of attributes of R that are not attributes of S .
- The result of DIVISION is a relation $T(Y)$. For a tuple t to appear in the result T , the values in t must appear in R in combination with every tuple in S .

Division (cont.)

To get $R \div S$, you could use:

- Project out Y, the attributes of R that are not in S

$$T_1 \leftarrow \pi_Y(R)$$

- T2 now contains all the tuples you do not want

$$T_2 \leftarrow \pi_Y((T_1 \times S) - R)$$

- Set difference removes these tuples, leaving just the tuples you do want

$$RESULT \leftarrow T_1 - T_2$$

R			
A1	A2	A3	A4
a	b	c	d
a	b	e	f
b	c	e	f
e	d	c	d
e	d	e	f
a	b	d	e

S	
A3	A4
c	d
e	f

R ÷ S	
A1	A2
a	b
e	d